

Power-of-Two (PoT) Weights in Large Language Models (LLMs)

Mahmoud Elgenedy

Computer Science, Stanford
melgened@stanford.edu



Motivation and Proposed Solution

Problem statement:

- In Large Language Models (LLMs), the number of model parameters has grown exponentially in the past few years, for example, from 1.5 billion parameters in GPT2 to 175 billion in GPT3.
- This raises a significant challenge for implementation, especially for Edge devices where memory and processing power are very limited.

Proposed Solution:

- We investigate reducing LLM complexity with special type of **quantization, power of two (PoT)**, for linear layers **weights** for both multi-layer perceptron **MLP** and **multi-head attention**.
- PoT not only provides **memory reduction** but more importantly provides significant computational reduction through converting **multiplication** to **bit-shifting**.
- PoT is investigated in few studies for CNN and general DNN [1]. Not avaiable clear studies for LLMs.

Summary of Results

For **124-M GPT-2** model, the PoT quantization results are shown to be very promising,

- Cross entropy loss degradation \approx [1.3-0.88] with number of bits range [4-6] to represent power levels.
- This results in **memory** reduction with up to a **factor of 8** and **processing power** reduction with a **factor of \approx 5**.

LLM Model and Datasets

NanoGPT [2] is a simple repository for training and tuning Generative Pre-trained Transformer GPT. The design follows GPT-2 model and can reproduce GPT-2 on OpenWebText dataset [3].

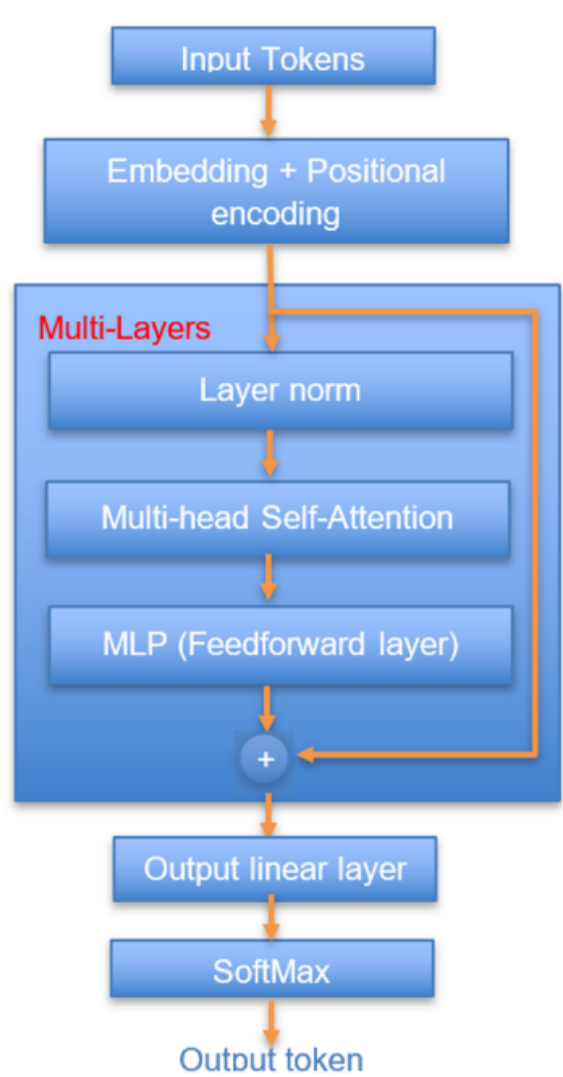
1. Shakespeare dataset

A small dataset (\approx 1 MB) consists of texts from William Shakespeare's plays. Our preliminary trials considering a character-level small LLM model trained only on Shakespeare dataset. We partitioned the data as 80%-10%-10% for training-val-testing. Character-level Tokenizer, from nano-GPT, is used to encode/decode dataset.

2. OpenWebText dataset

An open-source recreation of OpenAI's WebText dataset, which mainly used to train GPT-2. It is relatively large (\approx 54 GB) dataset consists of news articles, blog posts and stories. We evaluated the performance of quantization on OpenWebText dataset using a pretrained GPT-2 model [4]. The evaluation dataset is 0.05% of the total dataset. A Tokenizer from "tiktoken" by OpenAI is used.

Figure 1. LLM Model illustration and parameters



Parameter	Character-level	GPT- 2
Block Size (Context size)	64	1024
Number of Layers	4	12
Number of heads	4	12
Embedding Dimension	128	768
Dropout	0.2	0.2
Batch size	12	12
Vocab size	65	50257
Headen Layer dimension	4x128	4x768

Model Parameters

Method

Quantization: numbers can be represented in lower precision to help reduce both memory and processing requirements [5].

- Different types of quantization including symmetric, asymmetric quantization and Power-of-Two (PoT) (Table 1).
- Restricting weights to Power-of-Two (PoT) can significantly reduce processing power as multiplication is converted into bit-shifting.
- PyTorch provides different quantization frameworks[6], including Eager Mode, FX Graph Mode, and PyTorch 2 Export Mode.
- Both symmetric and Asymmetric are supported in PyTorch frameworks. However, PoT is not supported, and part of this work is to add it to PyTorch framework.
- We choose to use most recent **PyTorch 2 Export Mode** [7] as it is more flexible and easily scalable.

Table 1. Types of Quantization

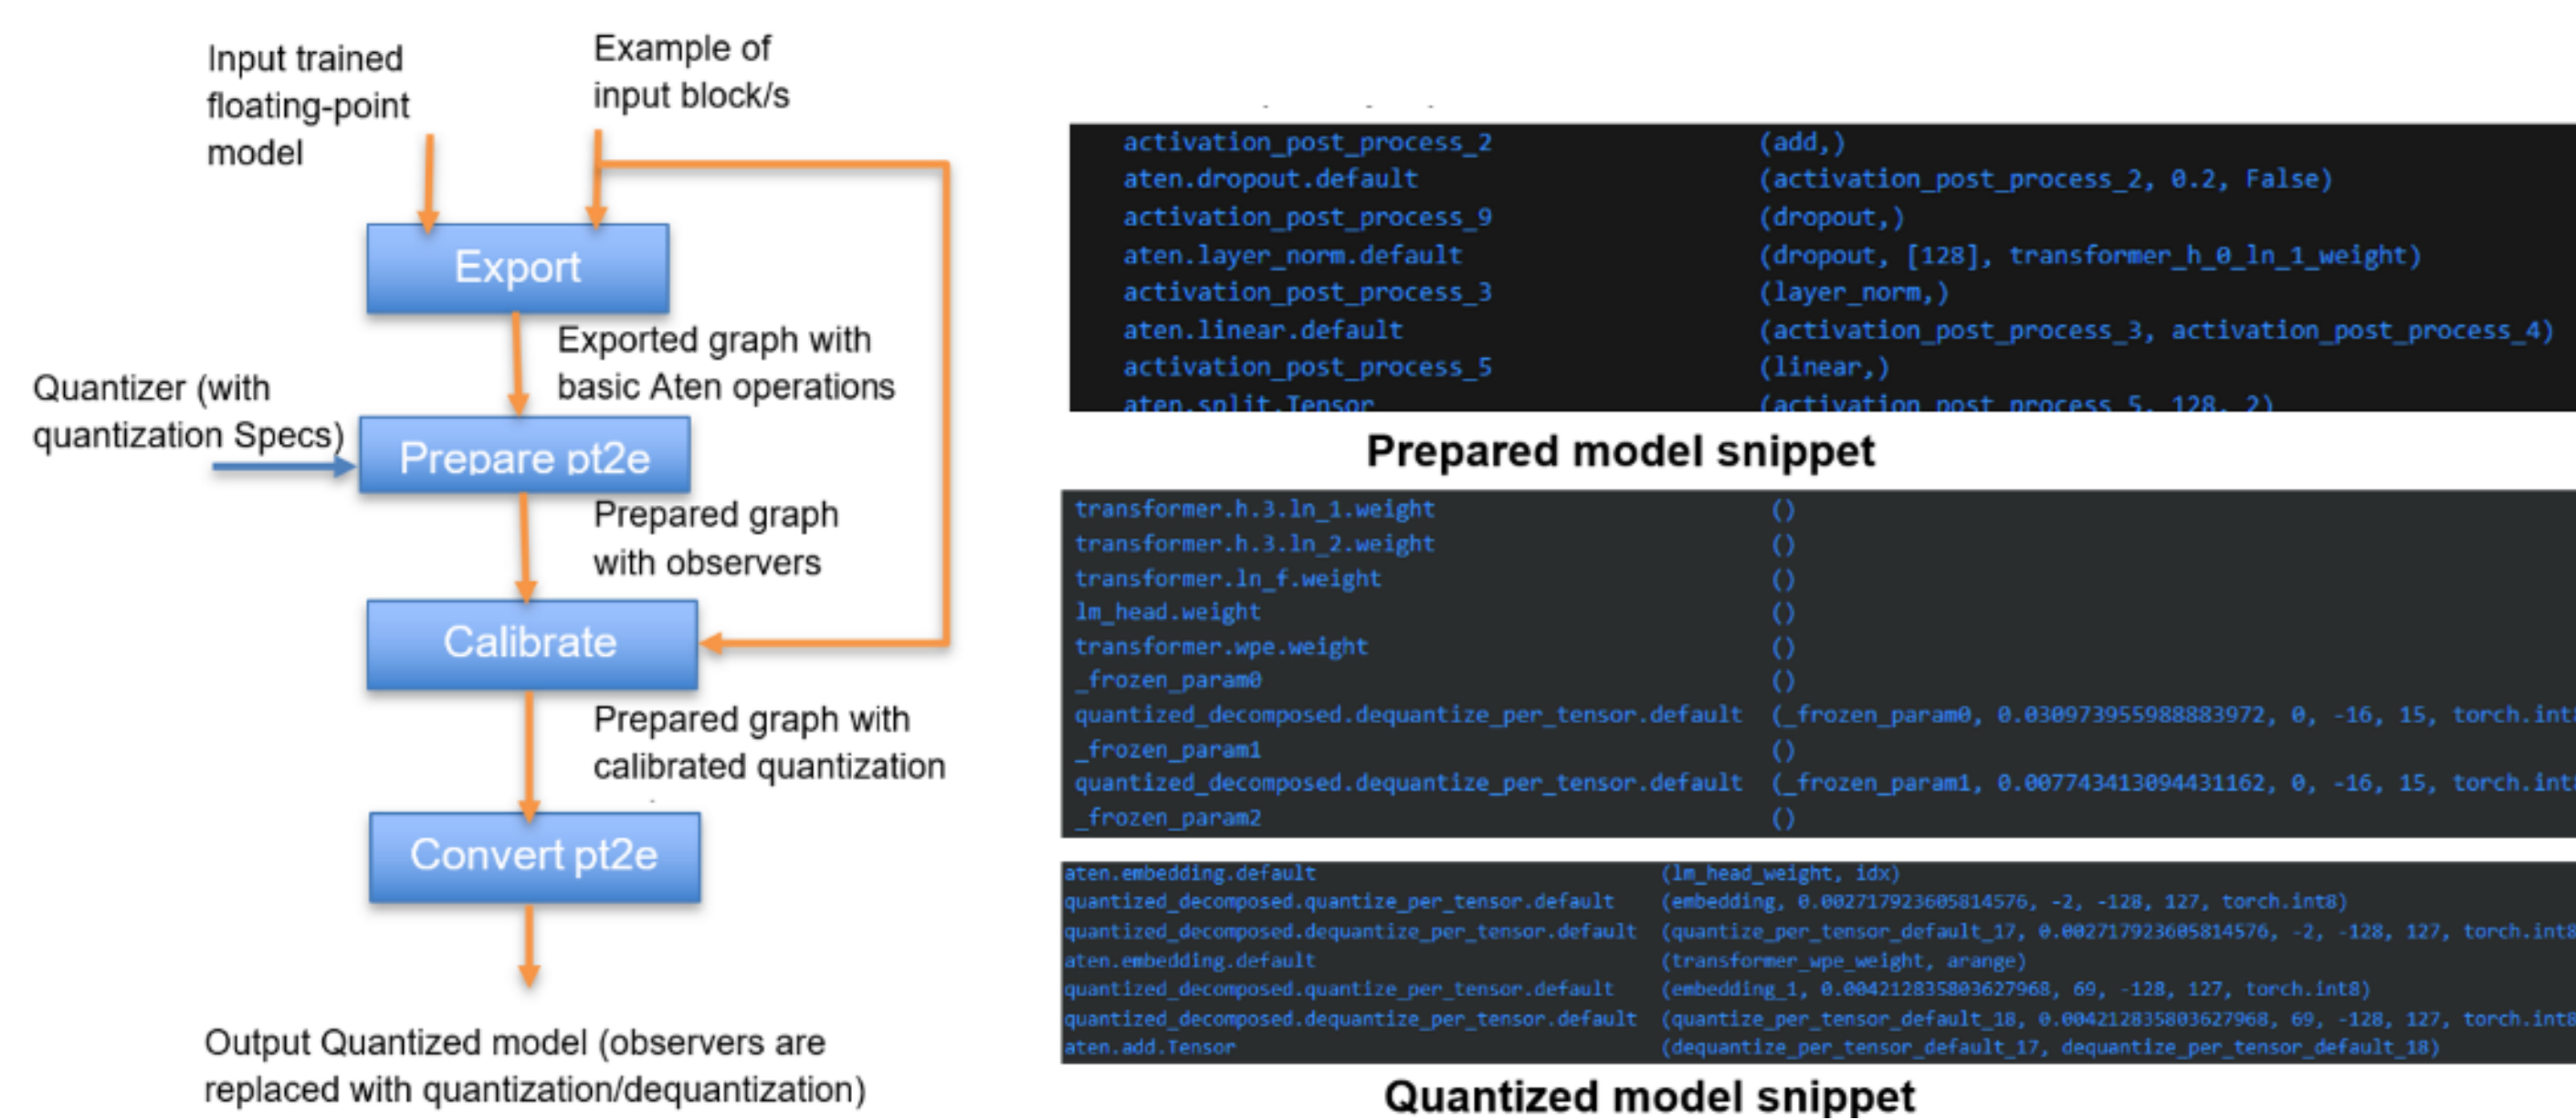
Quantization type	Conversion equation (input x, and output y)	Use case
Symmetric	$y = \text{round}(x / \text{scale})$	Weights
Asymmetric (Affine)	$y = \text{round}(x / \text{scale}) + \text{zp}$	Activations
Power of Two (PoT)	$y = 2^{\text{clip}(\text{round}(\log_2(x / \text{scale})))}$	Suggested in this proposal for Weights in MLP and Transformer

Quantized Model

Pytorch Steps to obtain quantized model^a, using **Post Training Quantization (PTQ)** approach,

- Export:** the trained floating-point model is broken down into a graph basic tensor operations.
- Prepare:** observers modules are added between operations (to estimate input/output range).
- Calibrate:** run the prepared model with realistic data so that observers can do calibration.
- Convert:** each observer is converted into quantize/dequantize operation based on calibration.
- PoT:** restrict weights and transformer LUTs to power of two^b.

Figure 2. PyTorch Quantization Framework



^aThis is not fully quantized model but simulates with good accuracy quantization loss

^bThis is additional step not in originally supported in Pytorch but added as part of this work

Performance Metrics

We use **cross-entropy loss** and **perplexity** to measure the performance degradation due to quantization.

To measure cross-entropy, we average over multiple iterations of calling the model (both floating and quantized) with input batches captured from testing set.

The cross-entropy $l_{ce}(t, y)$ with input k -dimension logits $t \in \mathbb{R}^k$ and target class y is defined as follows,

$$l_{ce}(t, y) = -\log P(y; t) = -\log \left(\frac{\exp(t_y)}{\sum_s \exp(t_s)} \right)$$

, and perplexity is just the exponential of the cross entropy defined as $\exp(l_{ce})$

Main Results

Table 2. Performance summary of Normal quantization versus PoT

Integer range	Normal quantization			Power of Two (PoT)		
	No. of bits	l_{ce}	Perplexity	No. of bits	l_{ce}	Perplexity
$[-2^5, 2^5]$	6 bits	7.7	2225.86	4 bits	9.563	14158.84
$[-2^7, 2^7]$	8 bits	3.329	27.92	4 bits	4.5	90
$[-2^{15}, 2^{15}]$	16 bits	3.1888	24.25	5 bits	4.3	73.7
$[-2^{20}, 2^{20}]$	32 bits	3.1883	24.24	6 bits	4.23	68.7
$[-2^{30}, 2^{30}]$	32 bits	3.187	24.23	6 bits	4.08	59.14

Figure 3. Output generated text with 4 bits PoT

But the committee says that the blame is not solely on the WHO. It calls for steps to bolster international research and development on emerging diseases between and during outbreaks, including the development of invasive viral infections as well as influenza. The problem was to bring into the groups lower ranks of those groups were that countries of birth. It is common knowledge among some groups to develop the groups above the governments. Most of the countries above the nations of the peoples life forms of those nations development are of the forms of those nations, using physical development among those country organisations. And the way and the ways of the groups develop the groups growth. So the groups developing activity among the societies evolution

- The PoT quantized model has very promising performance even at very low number of bits, **4-bits** (15 levels $[2^{-7}, 2^7]$). This results in huge reduction of memory requirements (a factor of **32/4=8**).
- Interestingly, at same number of bits (4-bits), uniform quantization shows a big failure (cross entropy \approx 7.7) and completely off generated text, while PoT shows only \approx 4.5 cross entropy loss, with very good quality output text shown in figure 3. This can be expected as PoT is **non-uniform** (log) quantization which can fit better the data with more levels for smaller values and less levels for outliers.
- Finally, PoT reduces all linear and transformer operations by factor of \approx 5 since multiplication requires approximately 5 clock cycles while bit shifting is just one clock cycle.

Future work

- Try **Quantization Assisted Training (QAT)** instead of Post Training Quantization (PTQ), as training quantized model can significantly enhance the accuracy of quantized model.
- Try relatively bigger models recently introduced to edge devices, e.g., **Llama** lightweight (1B).

References

- D. Przewlocka-Rus, S. S. Sarwar, H. E. Sumbul, Y. Li, and B. De Salvo, "Power-of-two quantization for low bandwidth and hardware compliant neural networks," *arXiv preprint arXiv:2203.05025*, 2022.
- Karpathy, "Karpathy-nanogpt: The simplest, fastest repository for training/finetuning medium-sized gpts."
- A. Gokaslan, "Dataset card for "openwebtext"."
- OpenAI-community, "Gpt-2 model."
- A. Gholami, S. Kim, Z. Dong, Z. Yao, M. W. Mahoney, and K. Keutzer, "A survey of quantization methods for efficient neural network inference," in *Low-power computer vision*, pp. 291–326, Chapman and Hall/CRC, 2022.
- "Quantization, quantization - pytorch 2.6 documentation."
- "Pytorch 2.0 export post training static quantization."